



Microservices & Apache Kafka®

Part 1 – The Data Dichotomy: Rethinking the way we treat data and services



Series Schedule

- **Session 1: The Data Dichotomy: Rethinking the way we treat data and services**
- Session 2: Building Event-Driven Services with Apache Kafka
- Session 3: Putting the 'Micro' into Microservices with Stateful Stream Processing

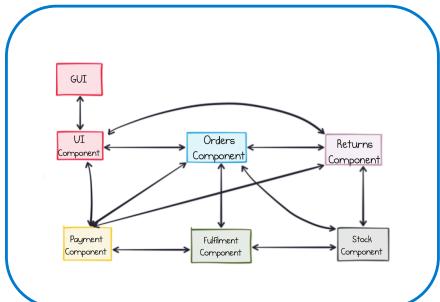


What are
microservices
really about?

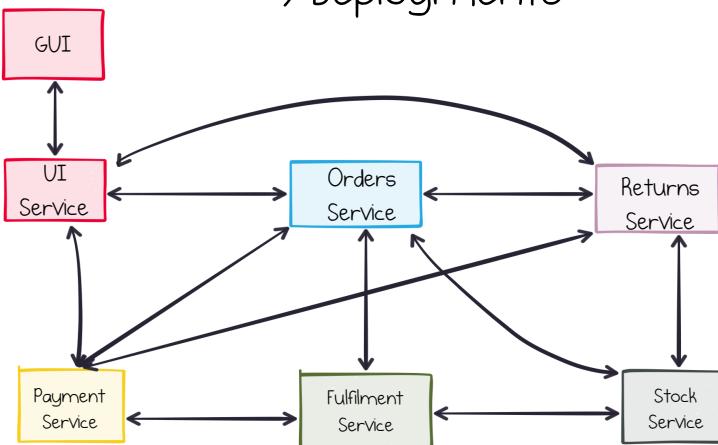


Splitting the Monolith?

Single Process
/Code base
/Deployment

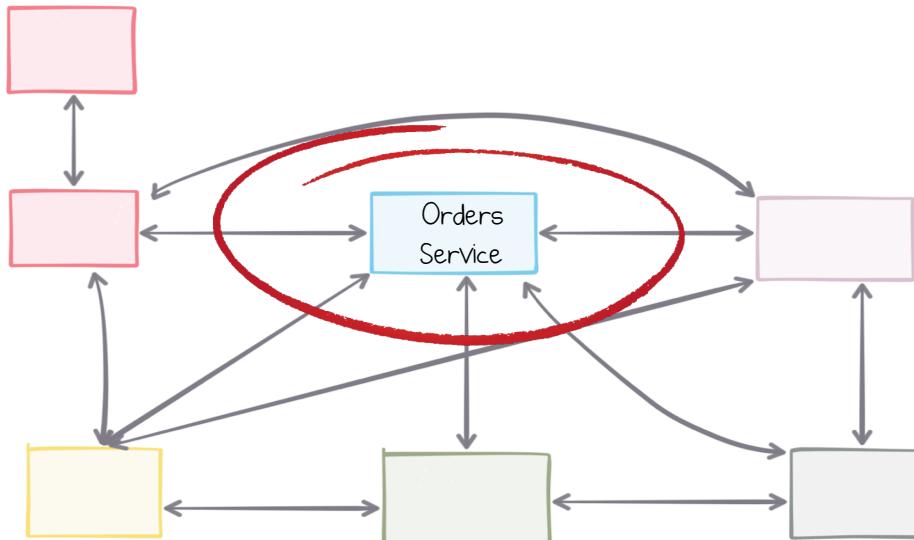


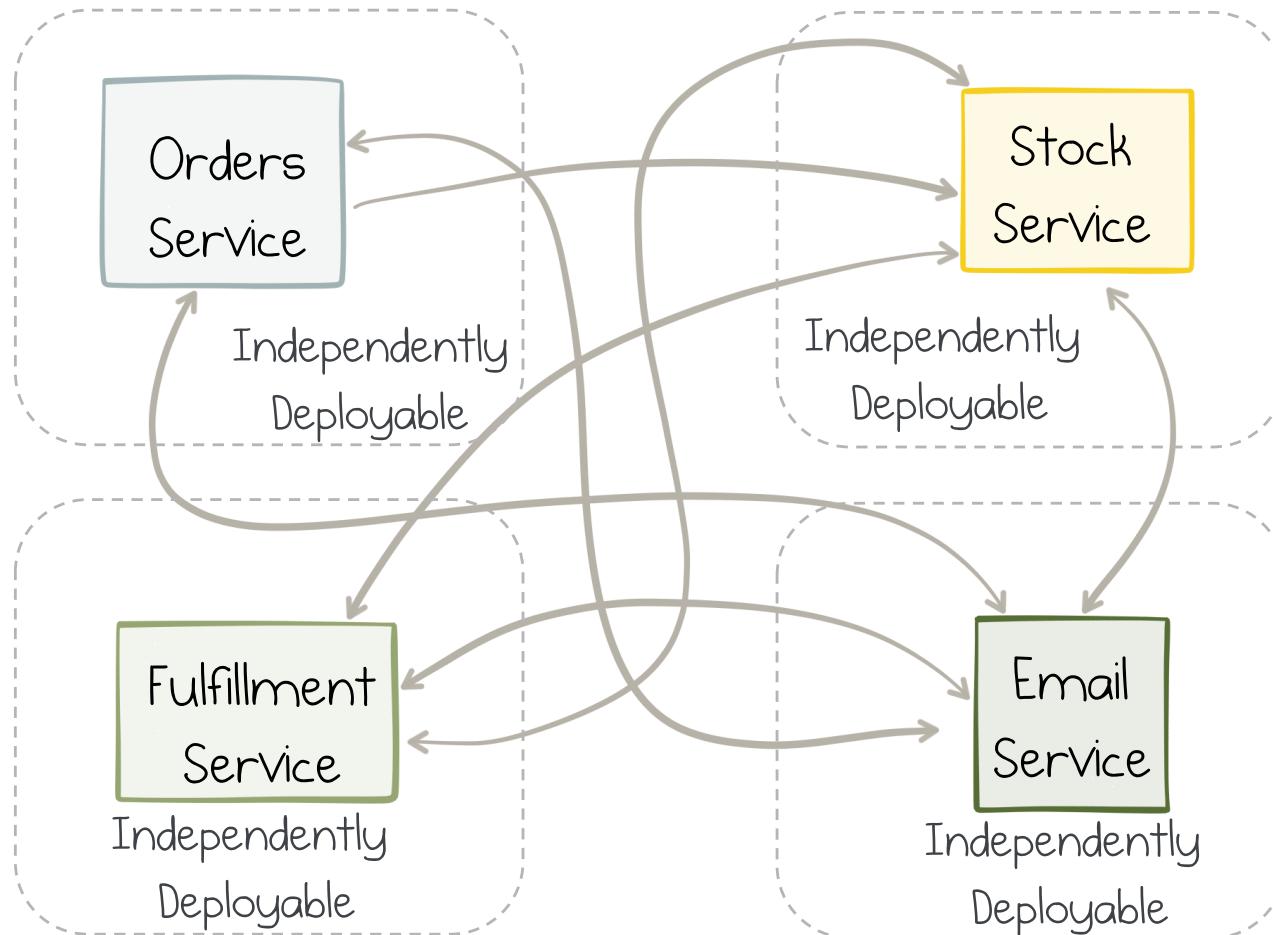
Many Processes
/Code bases
/Deployments





Autonomy?







Independence is
where services get
their value



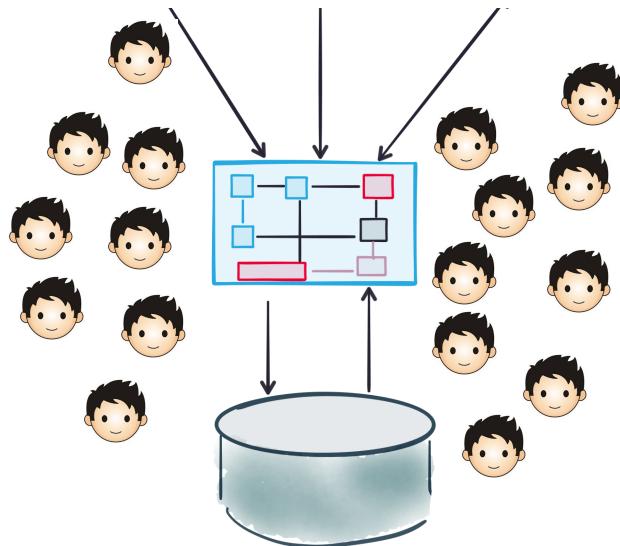
Allows Scaling



but not just in terms of
data/load/throughput...



Scaling in terms of people

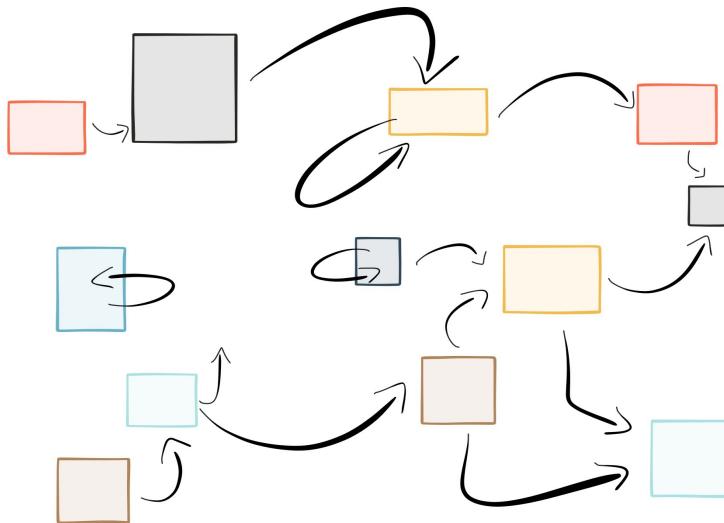


What happens when we grow?



Companies are inevitably a collection of applications

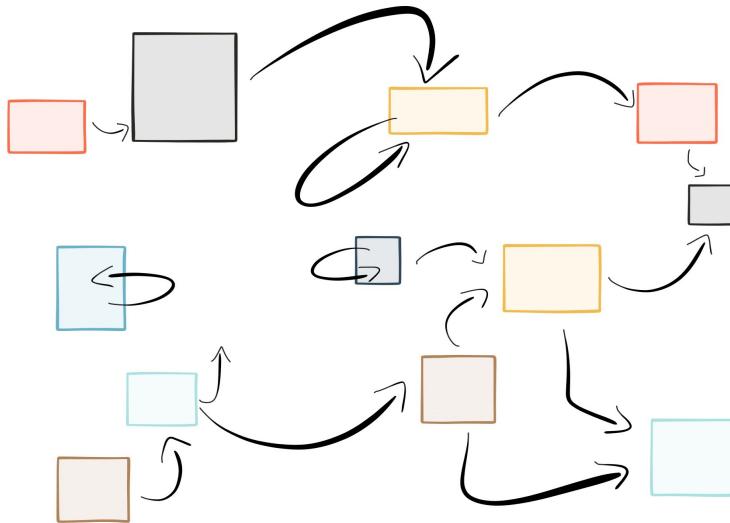
They must work together to some degree



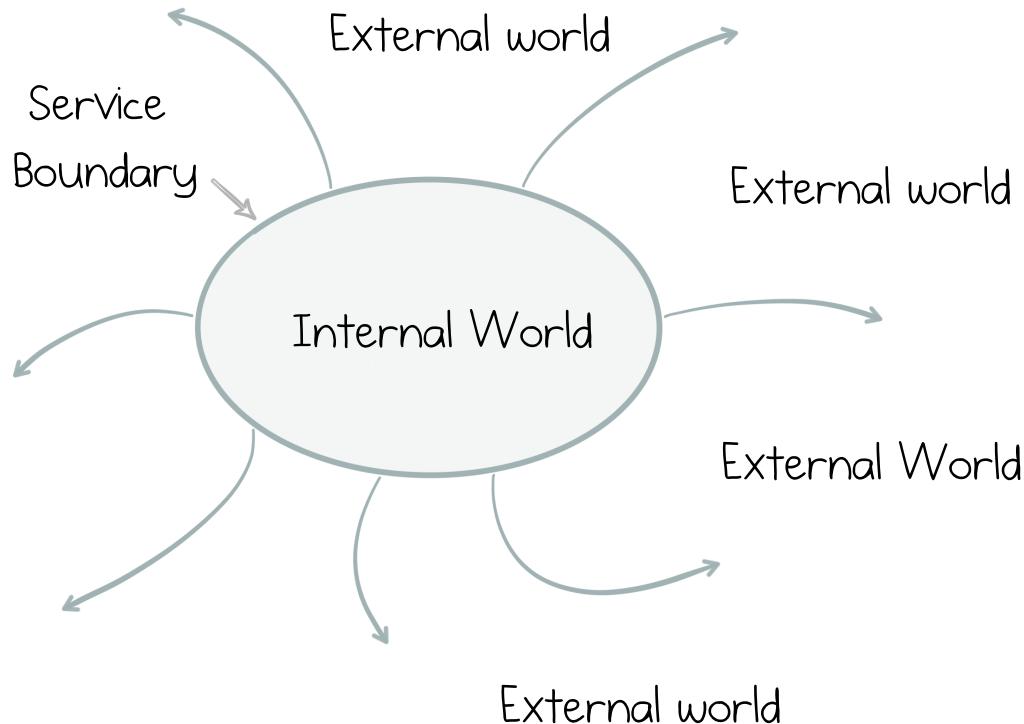


Interconnection is an afterthought

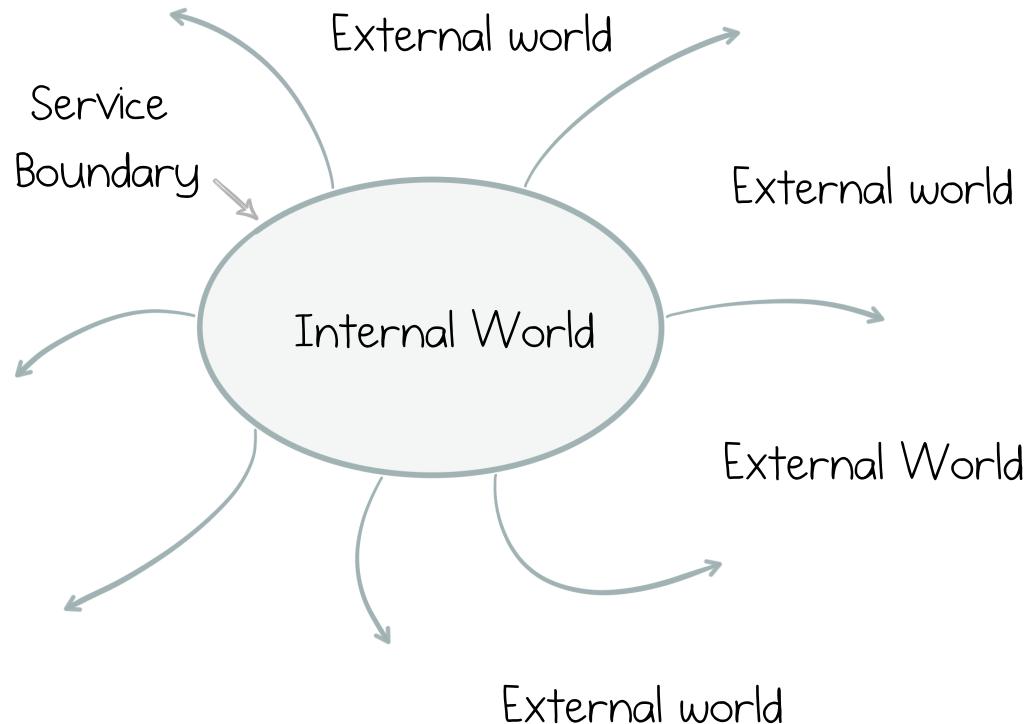
FTP / Enterprise Messaging etc



Services Force Us To Consider The External World



External World is something we should Design For



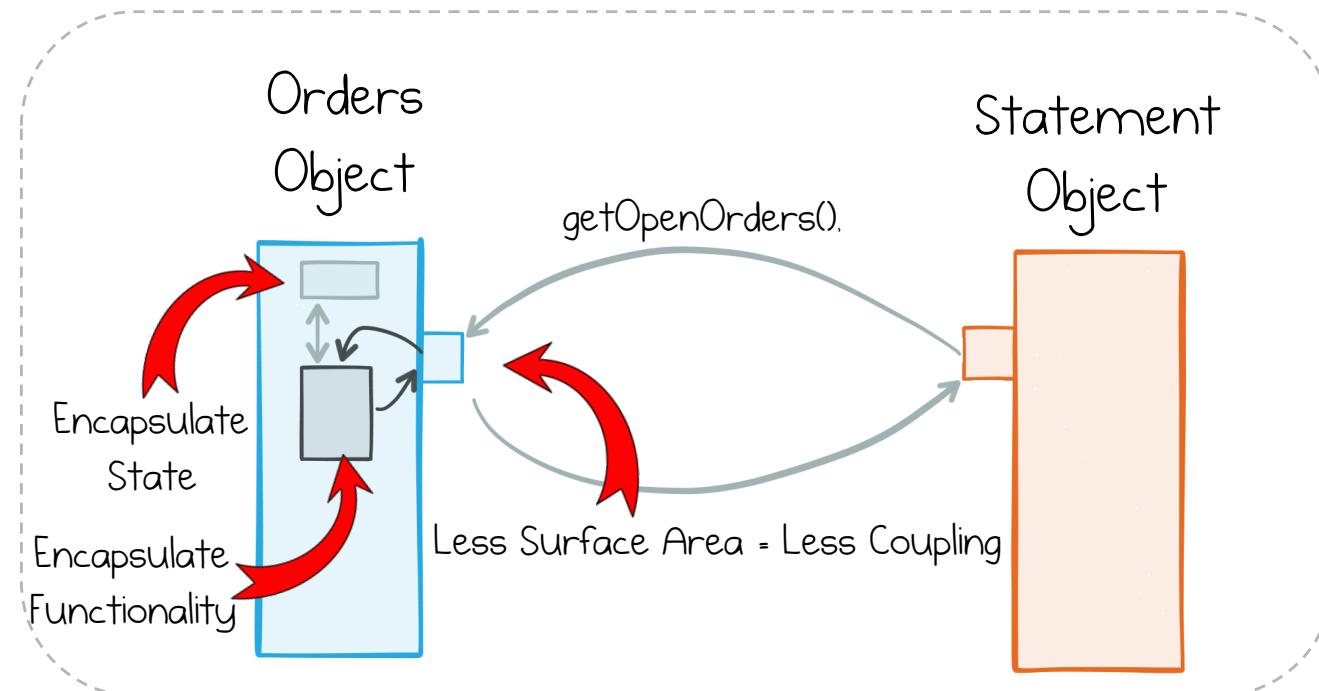


But this
independence
comes at a cost

\$\$\$



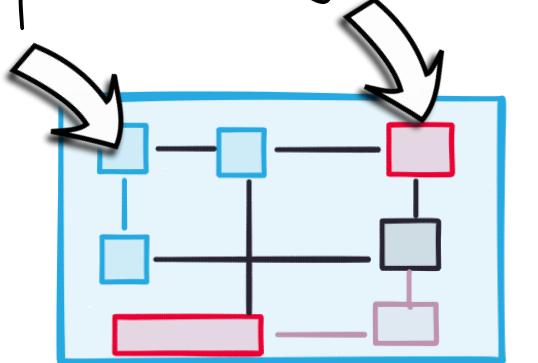
Consider Two Objects in one address space



Encapsulation \Rightarrow Loose Coupling



Change 1
Change 2



Redeploy

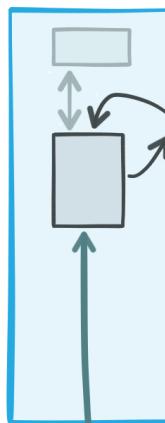


Singly-deployable
apps are easy



Independently Deployable

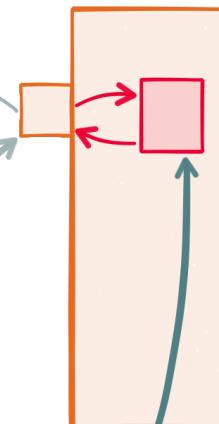
Orders Service



getOpenOrders()

Independently Deployable

Statement Service



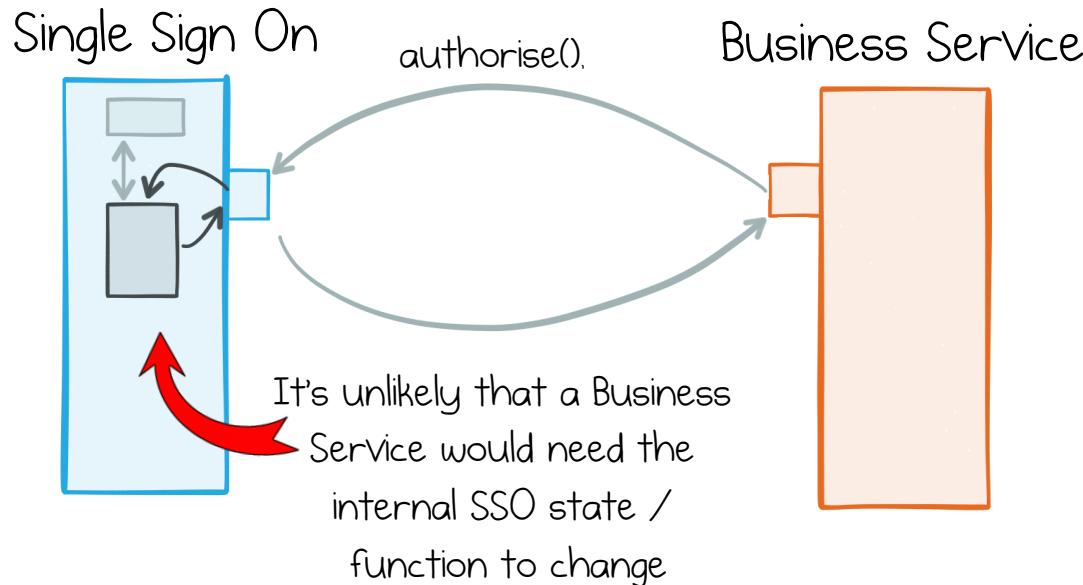
Synchronized
changes are painful



Services work best
where requirements
are isolated in a single
bounded context



Single Sign On





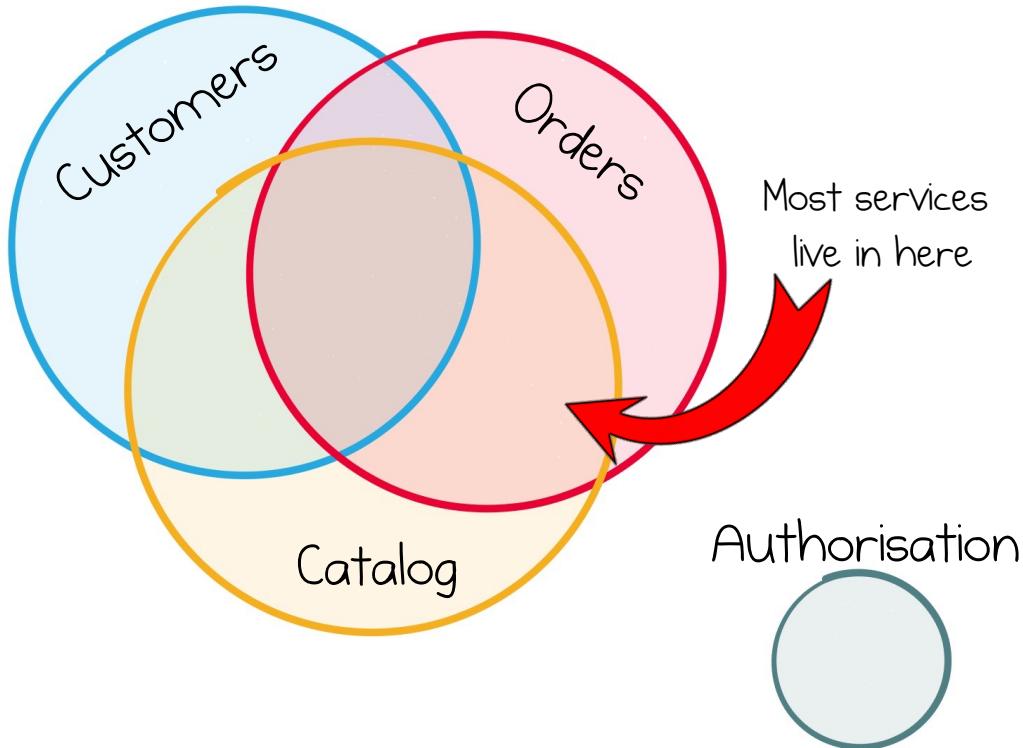
SSO has a tightly
bounded context



But business
services are
different



Most business services share the same core stream of facts.





The futures of
business services
are far more
tightly intertwined.



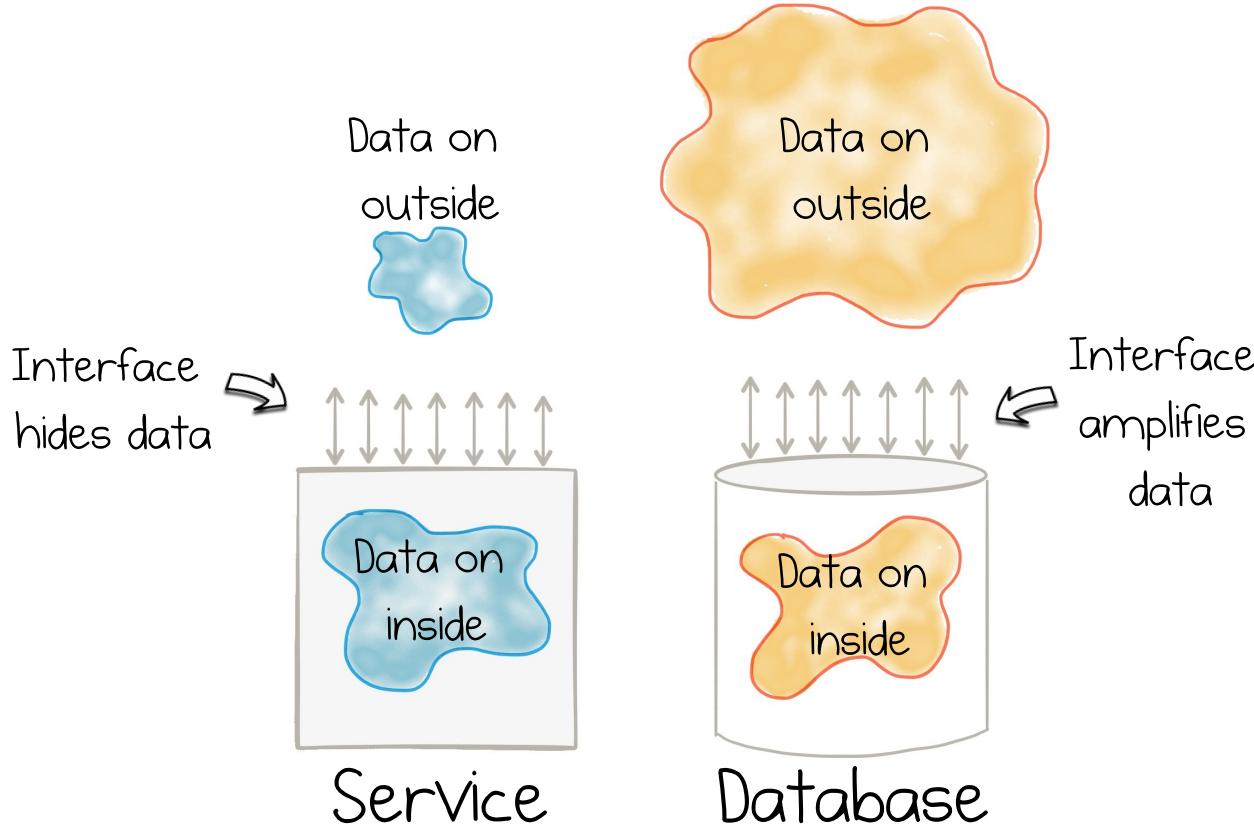
We need encapsulation to hide
internal state. Be loosely coupled.

But we need the freedom to slice
& dice shared data like any other
dataset



But data systems
have little to do
with encapsulation

Databases amplify the data they hold





THE DATA DICHOTOMY

Data systems are about exposing data.
Services are about hiding it.



This affects
services in one of
two ways

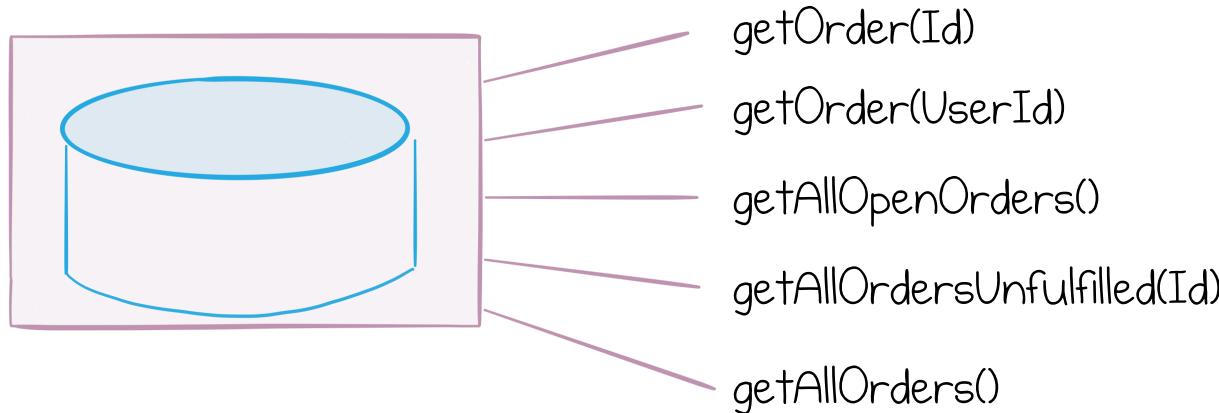


Either (1) we constantly add to the interface,
as datasets grow





(I) Services can end up looking like
kookie home-grown databases

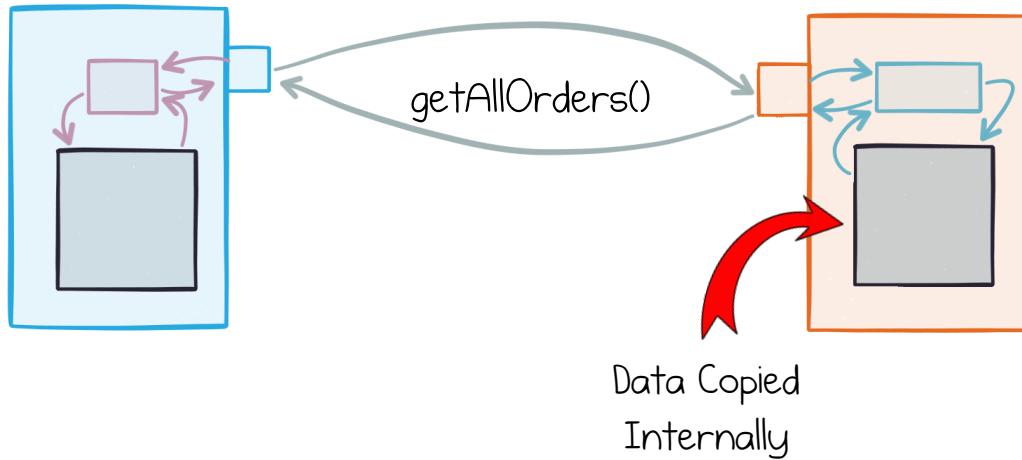




...and DATA amplifies
this service boundary
problem



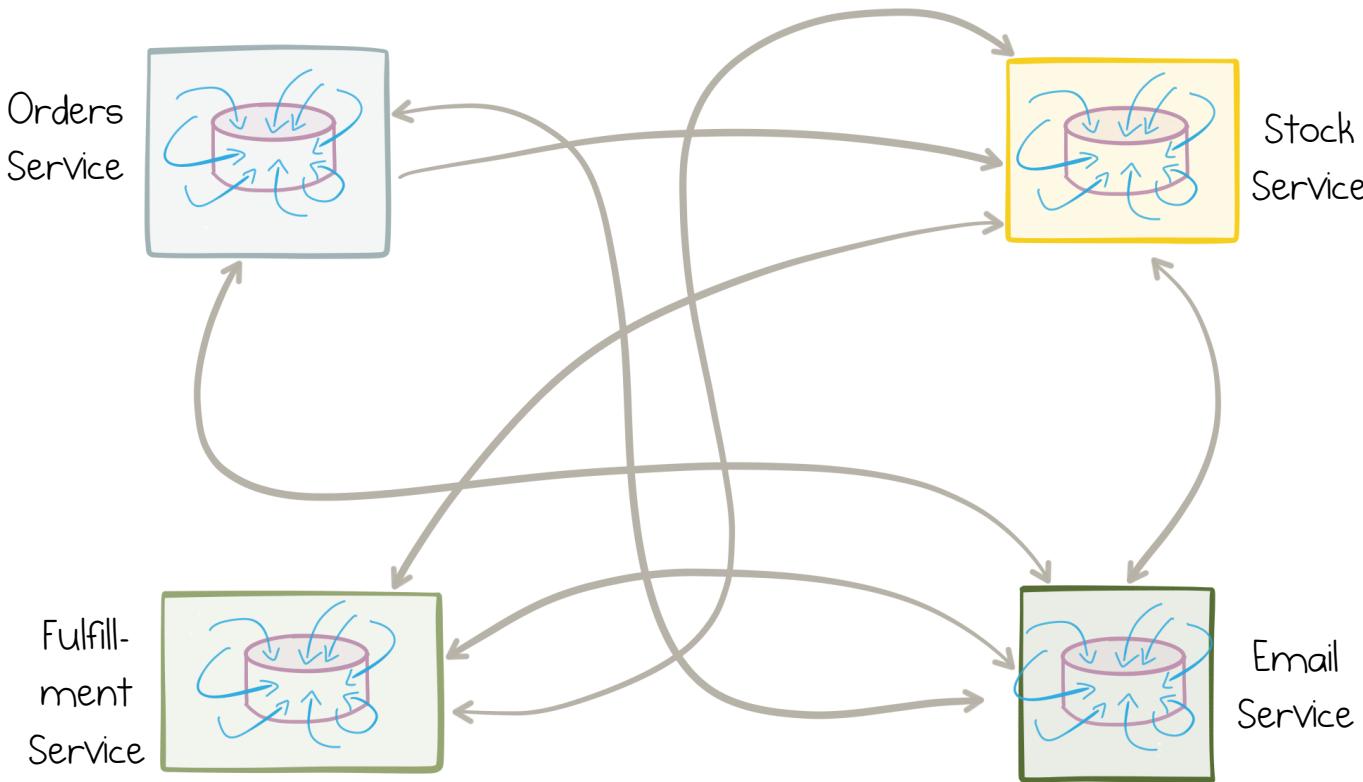
(2) Give up and move whole datasets en masse





This leads to a
different problem

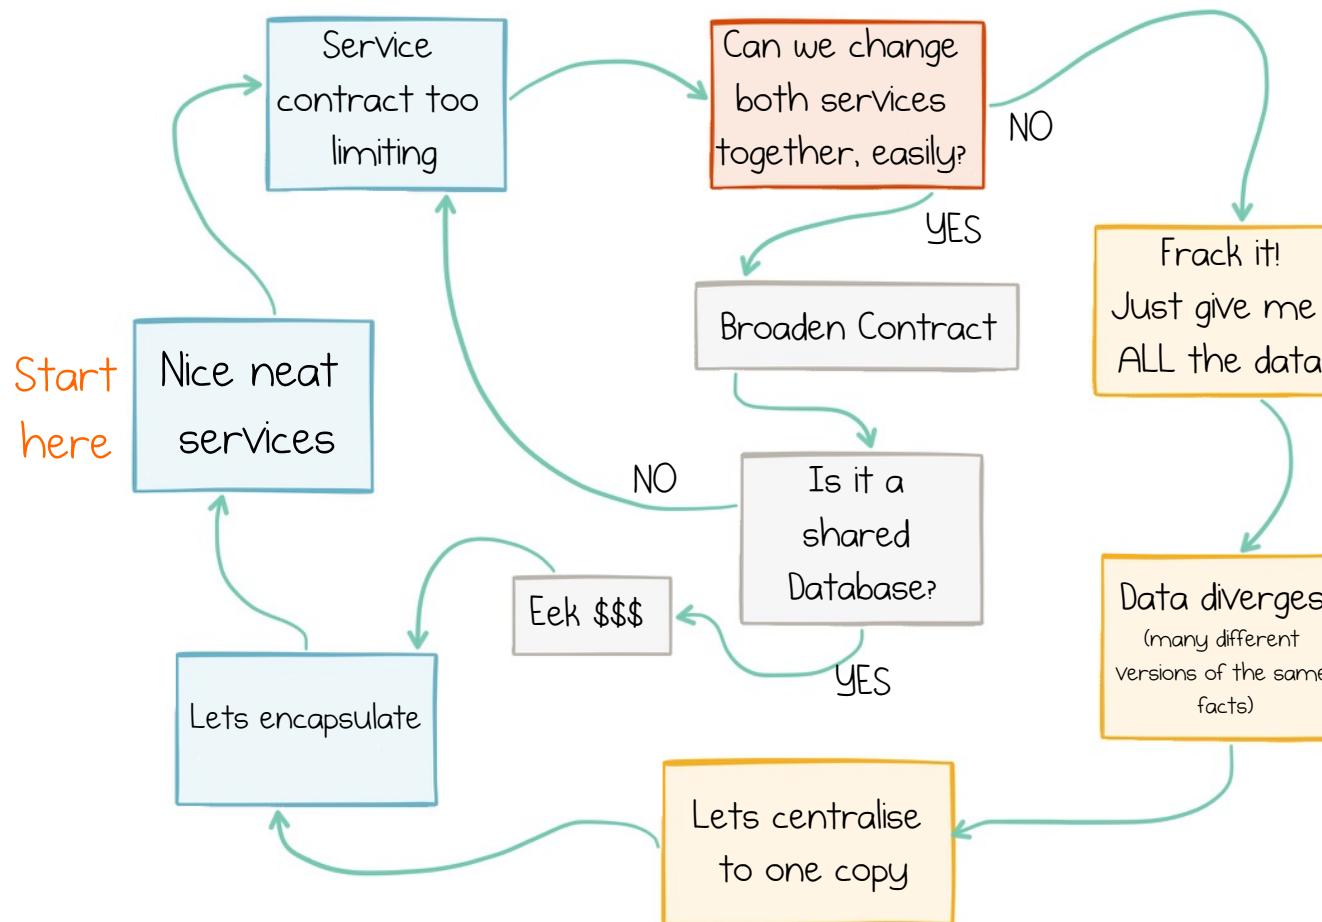
Data diverges over time





The more
mutable copies,
the more data will
diverge over time

Cycle of inadequacy:





These forces compete in
the systems we build

Encapsulation *vs* **Accessibility**



Accessibility *vs* **Divergence**



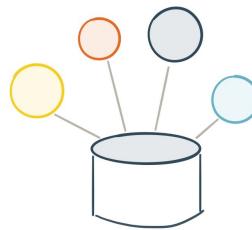
Divergence



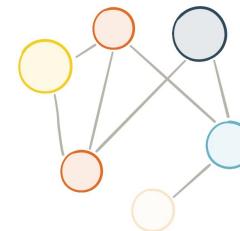


No perfect solution

Shared
database



Service
Interfaces

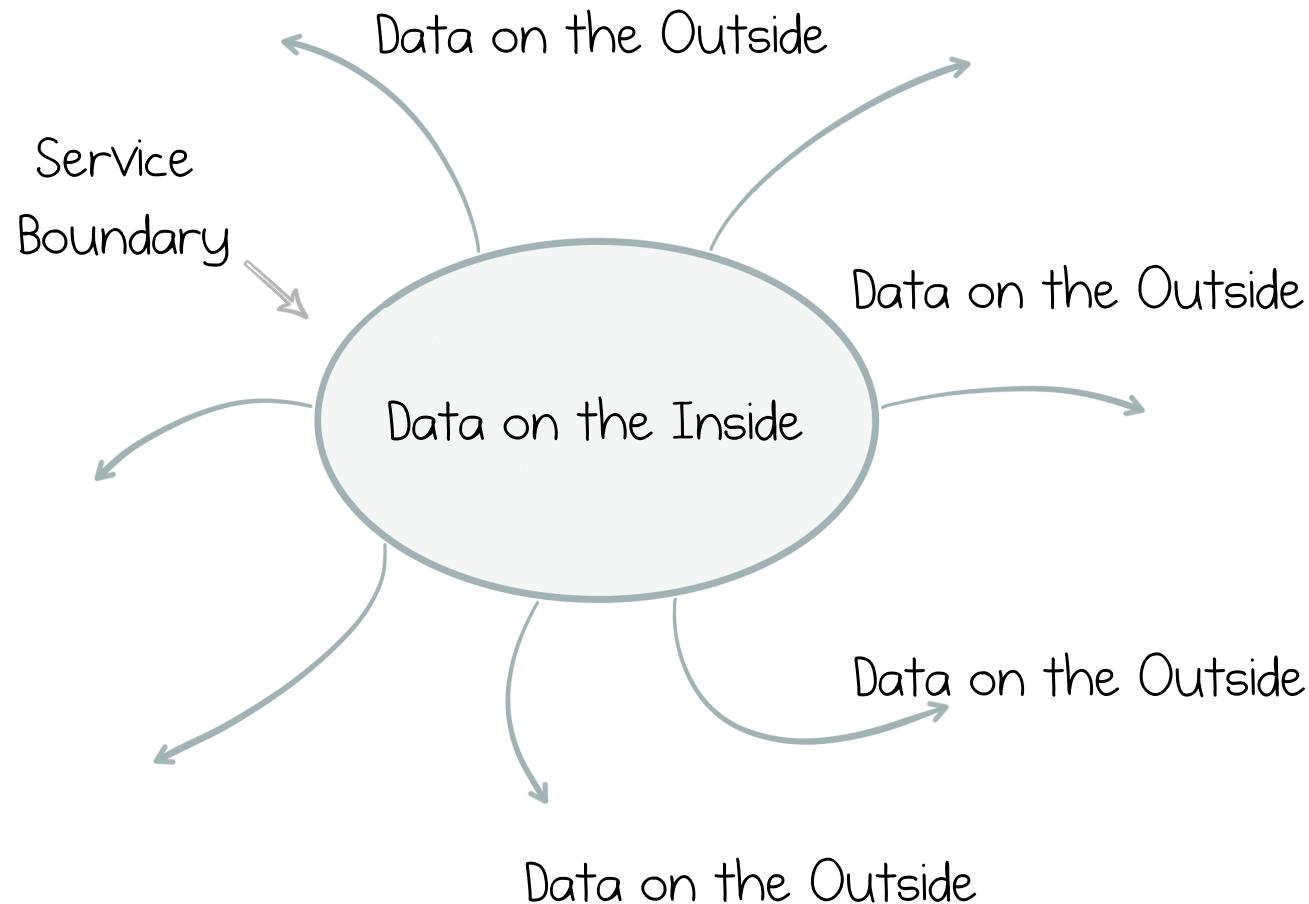


Better
Accessibility

Better
Independence



Is there a better
way?





Make
data-on-the-outside
a 1st class citizen



Separate Reads &
Writes with
Immutable
Streams

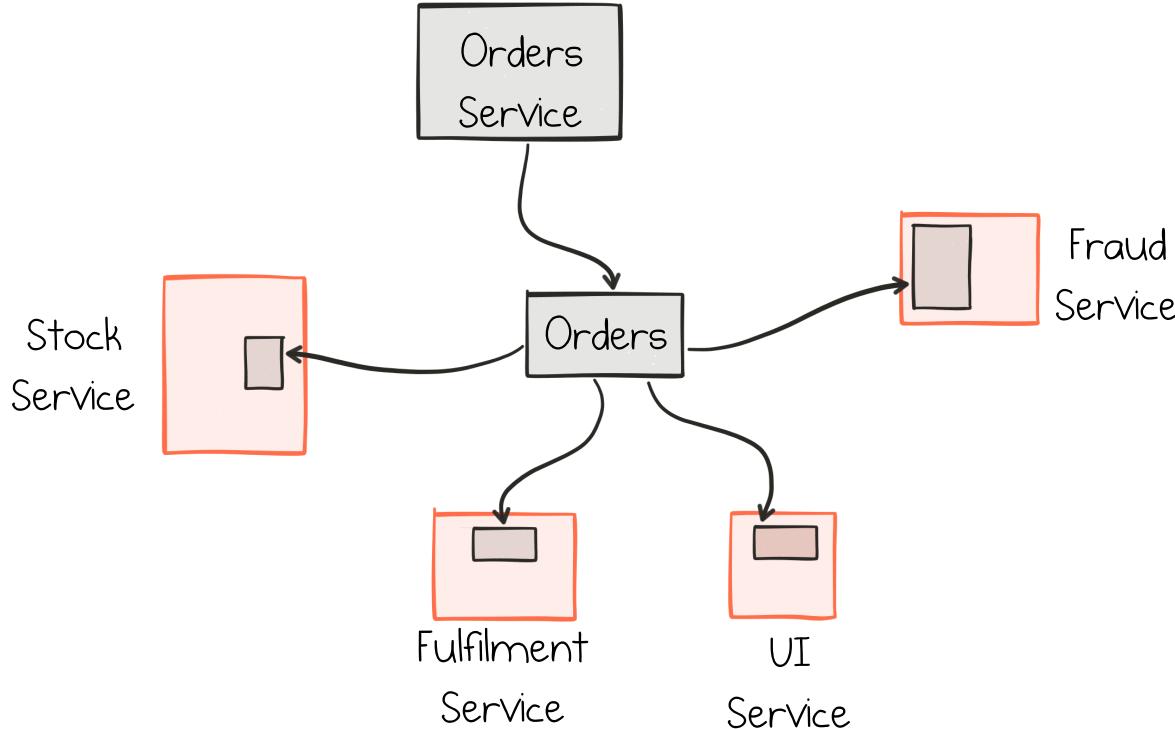


CQRS

- Writes => “Normalized” into each service
- Reads => “Denormalized” into Services
- One canonical set of streams

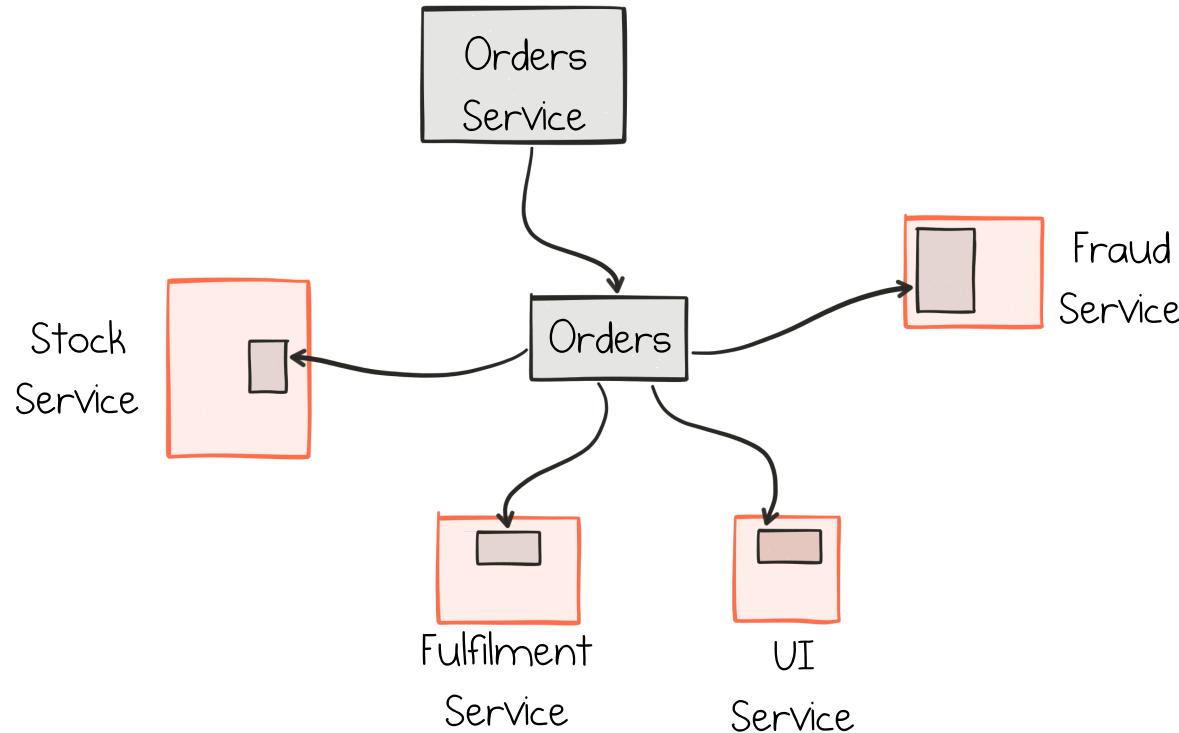


Writes enter via the Orders Service (which manages the workflow of orders)

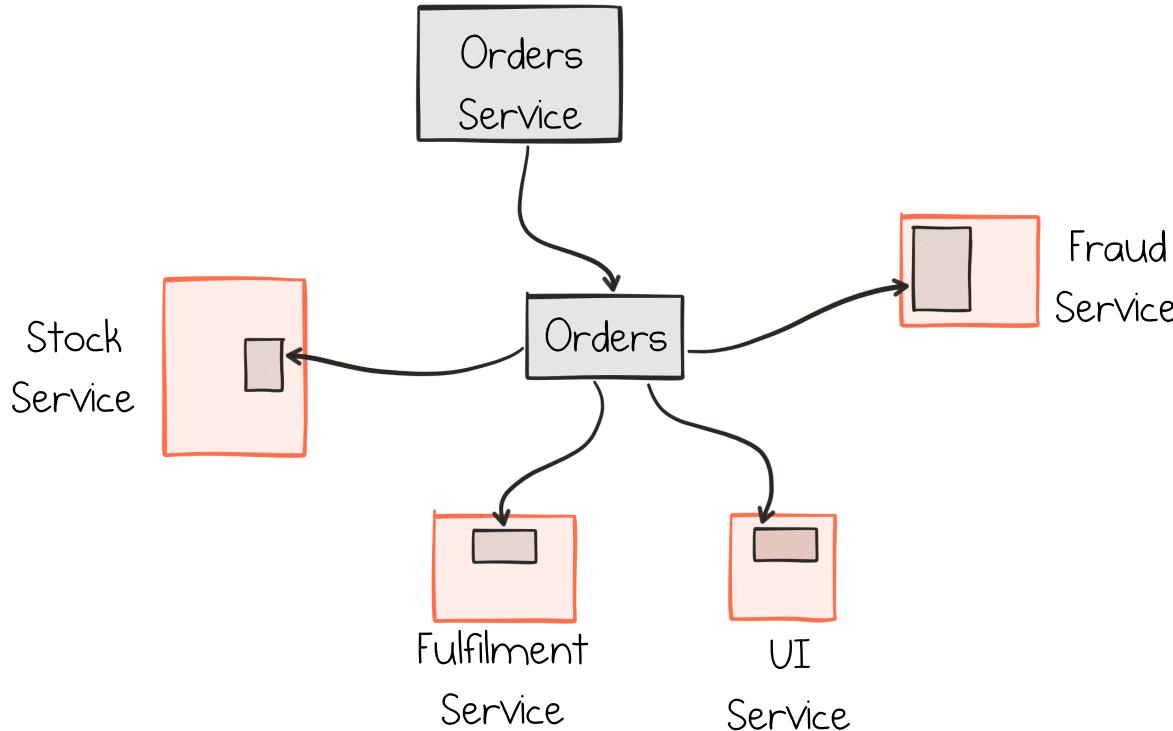




Readers use streams, materialized in each service



Important: data is not retained, it's just a view over those centralized streams



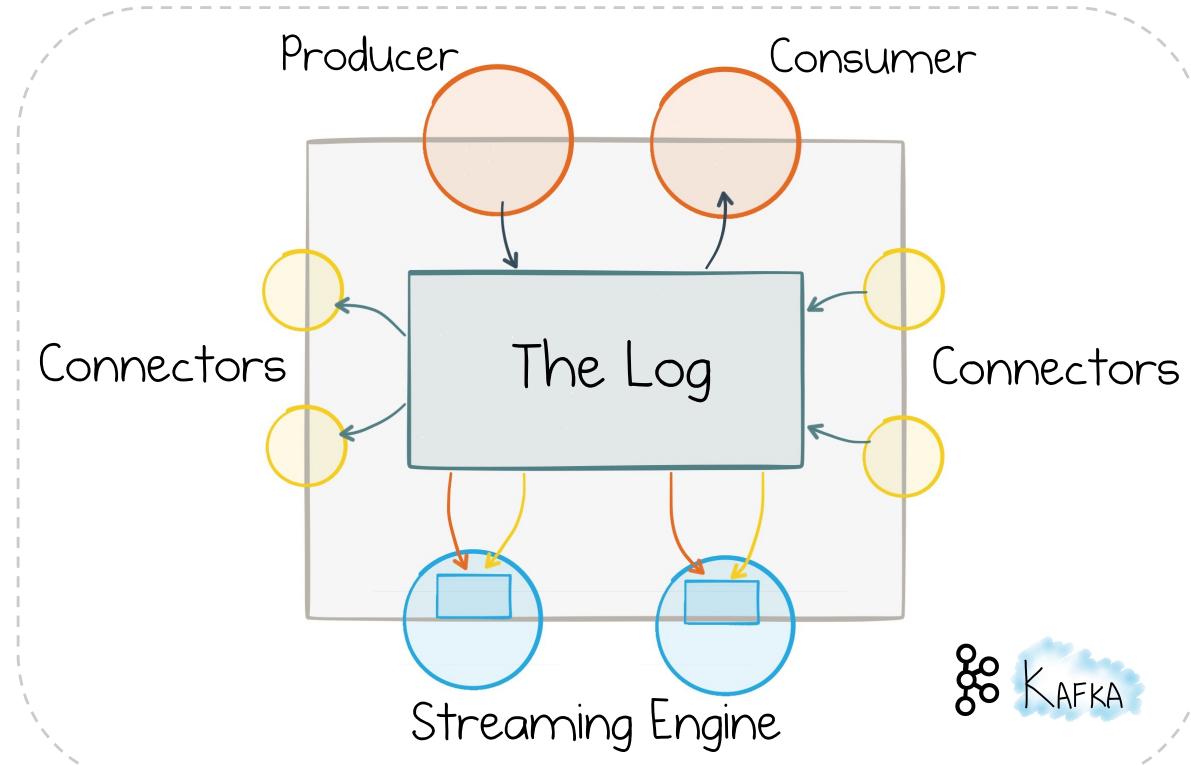


Use a Stream
Processing tool-kit



Kafka is a Streaming Platform

Kafka: a Streaming Platform

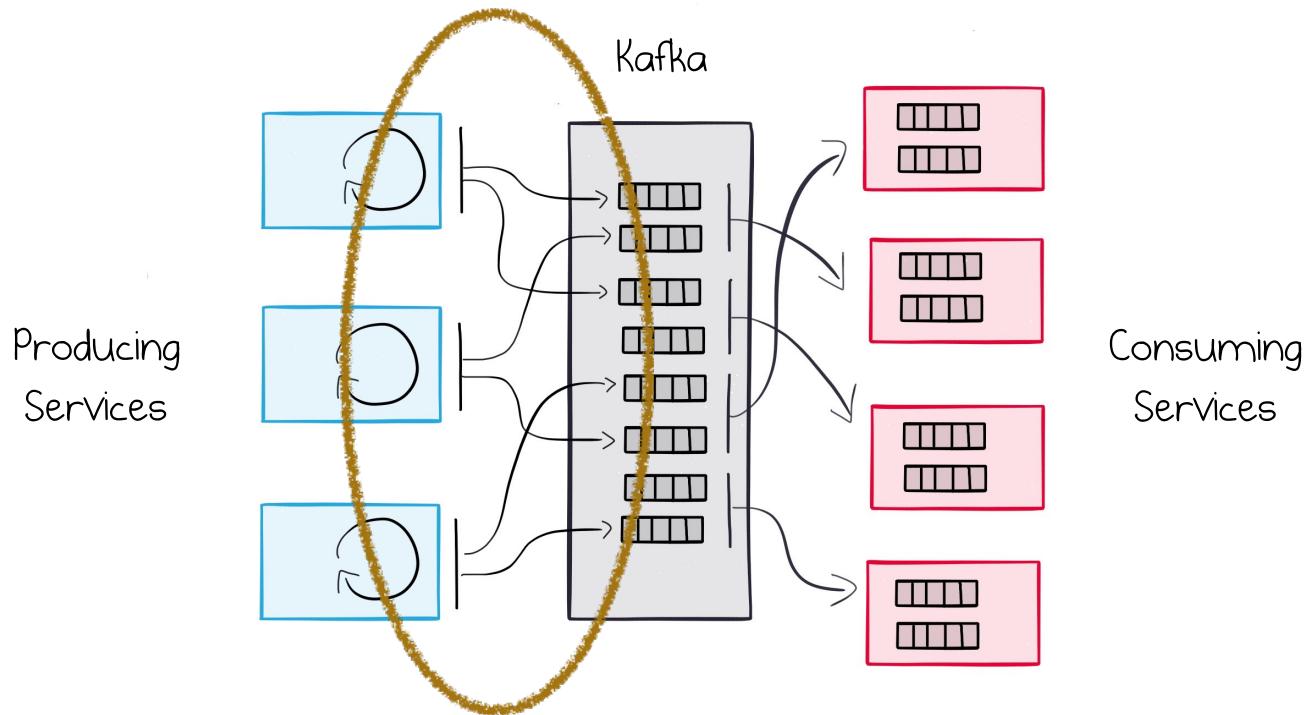




The Log

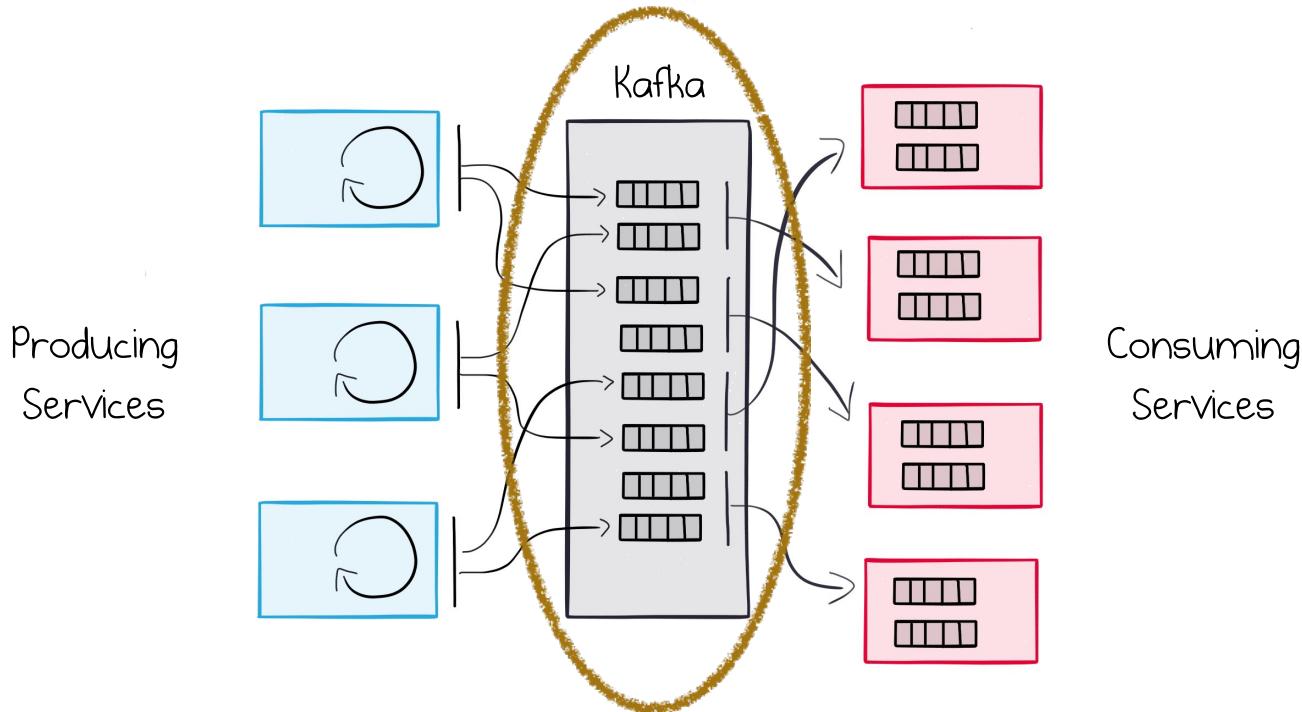


Shard on the way in



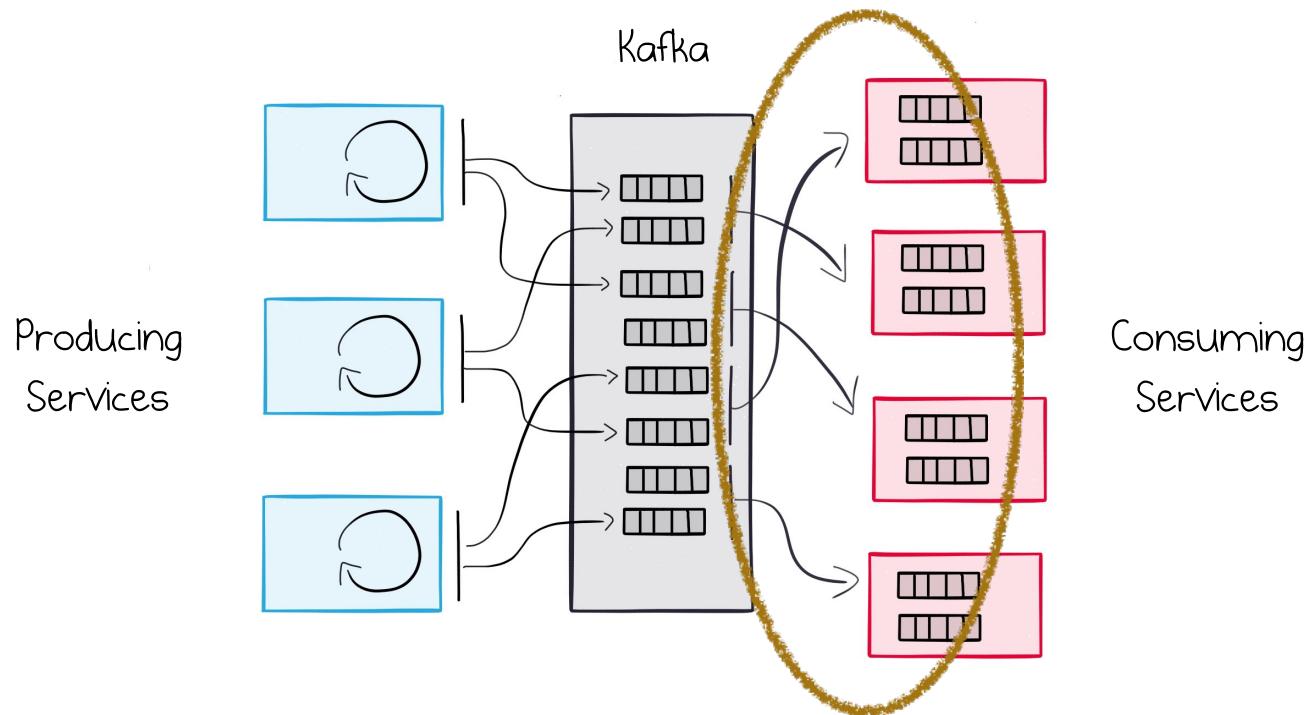


Each shard is a queue





Consumers share load

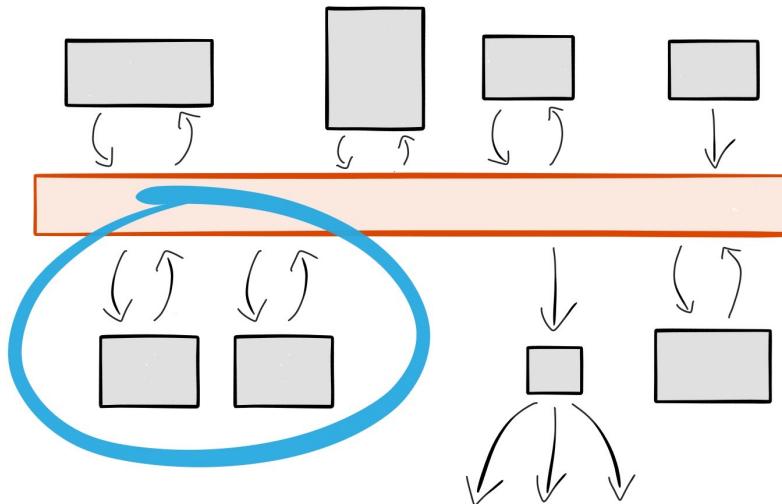




Scaling becomes a
concern of the
broker, not the
service

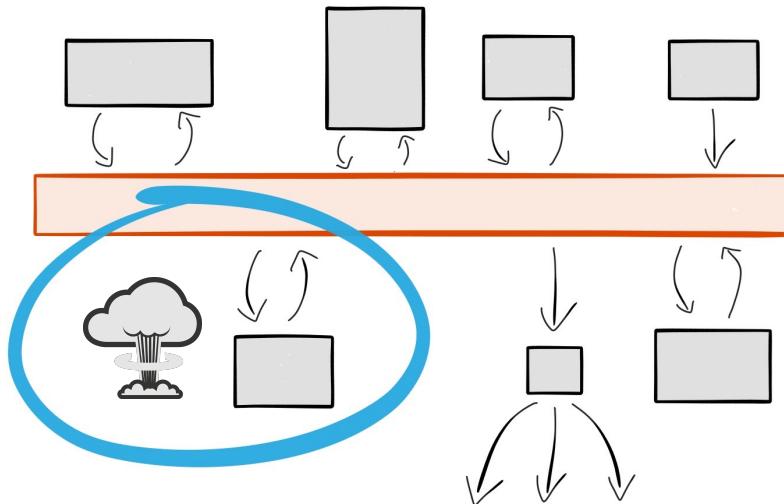


Load Balanced Services



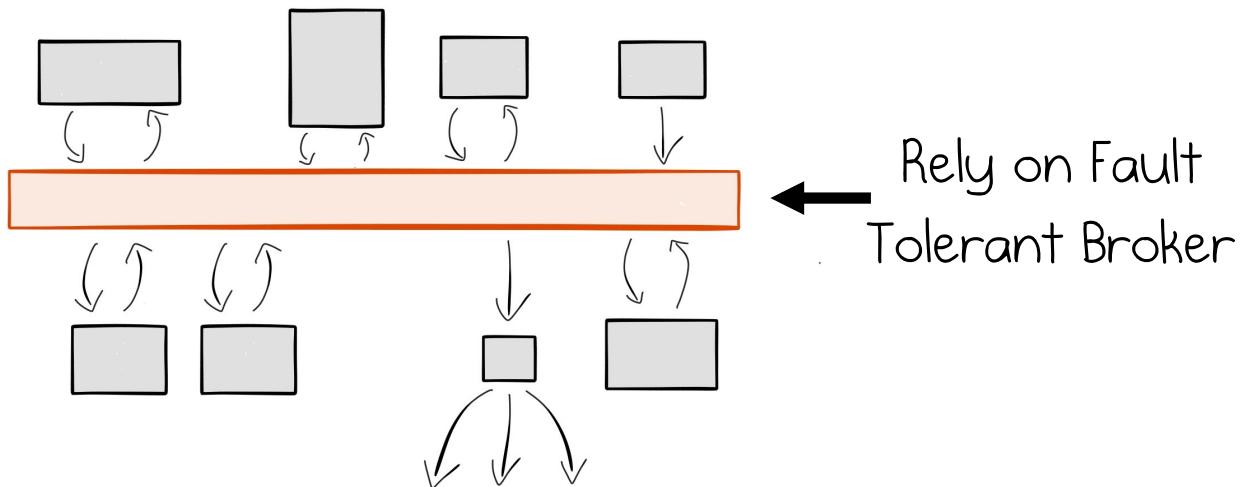


Fault Tolerant Services



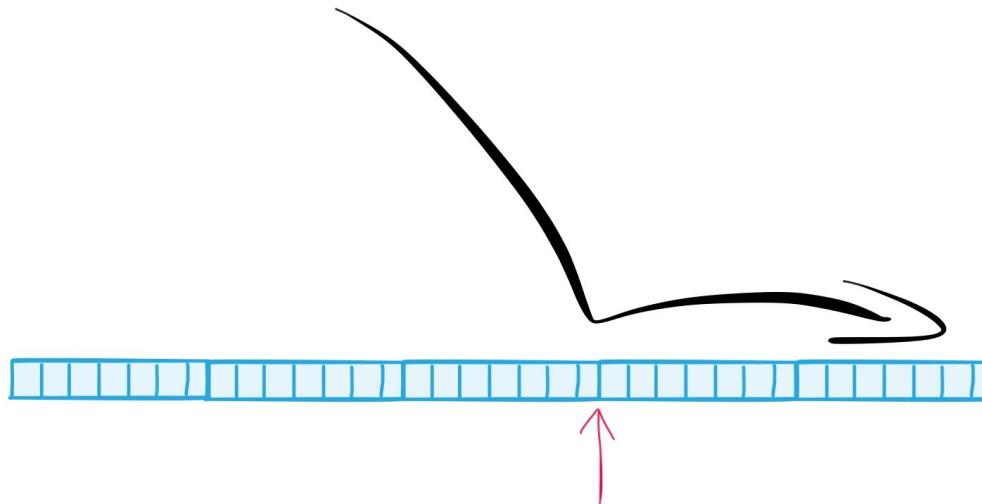


Build 'Always On' Services





Reset to any point in the shared
narrative

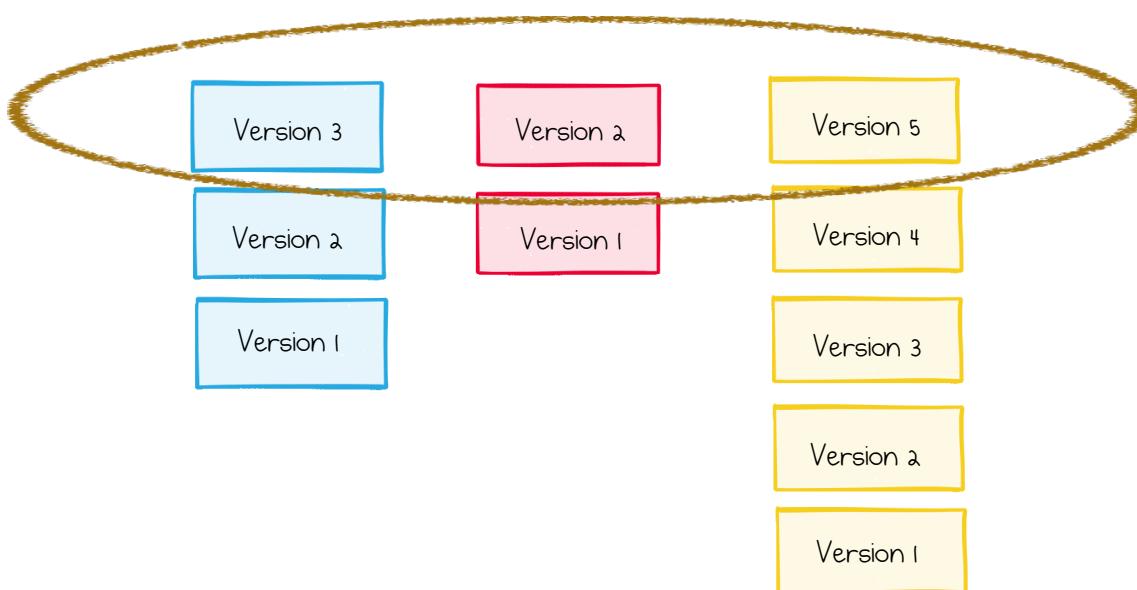


Rewind & Replay



Compacted Log

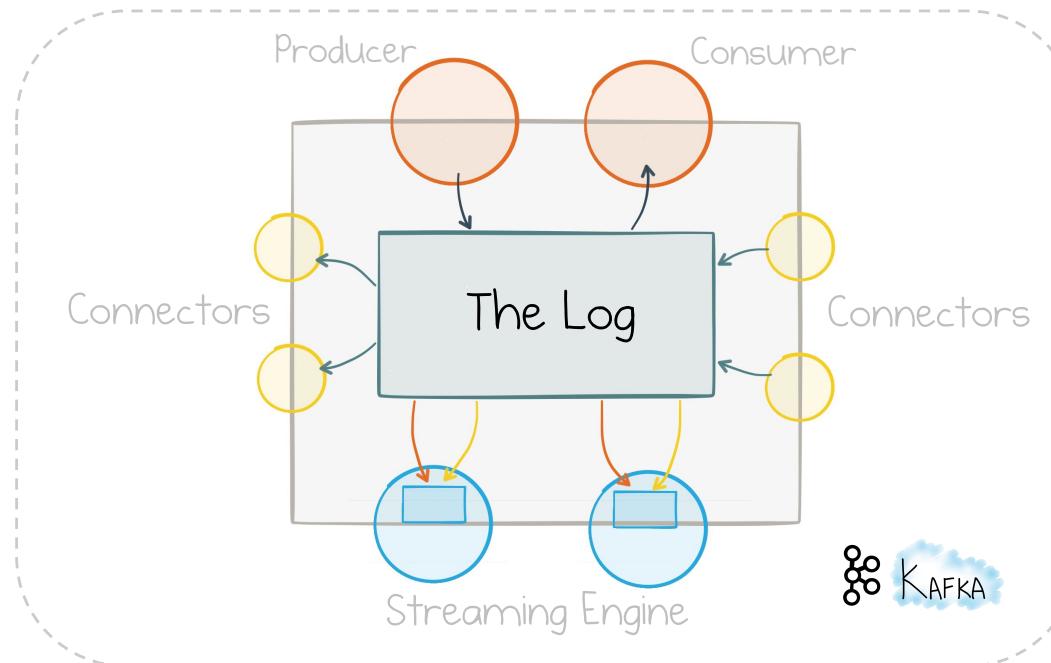
(retains only latest version)





Service Backbone

Scalable, Fault Tolerant, Concurrent, Strongly Ordered, Retentive





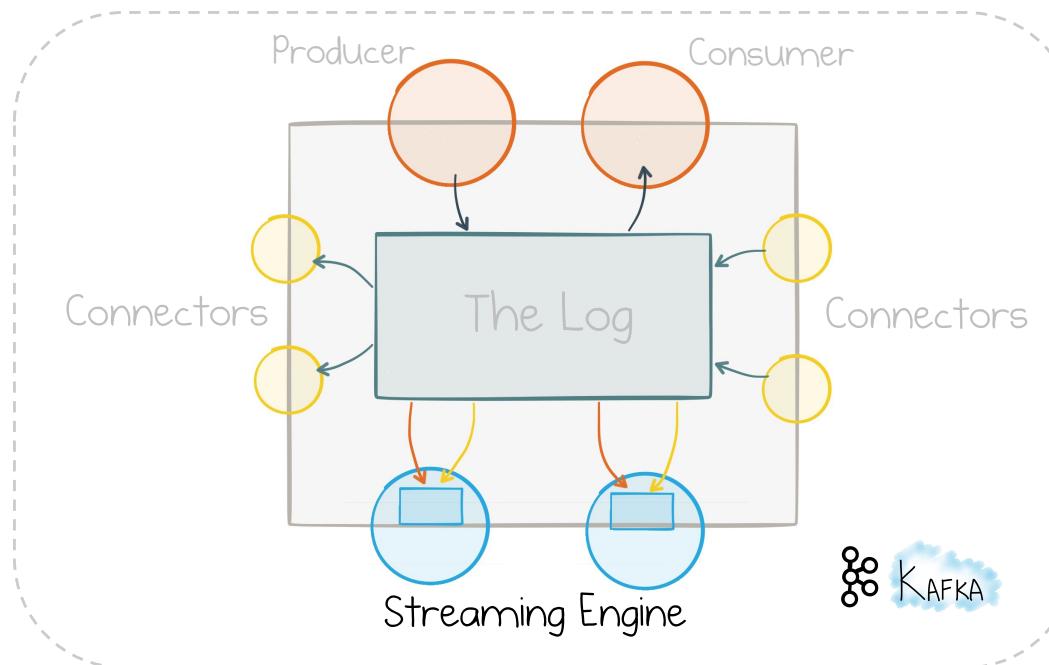
A place to keep the
data-on-the-outside
as a immutable
narrative



Now add Stream
Processing

Kafka's Streaming API

A general, embeddable streaming engine





What is Stream Processing?





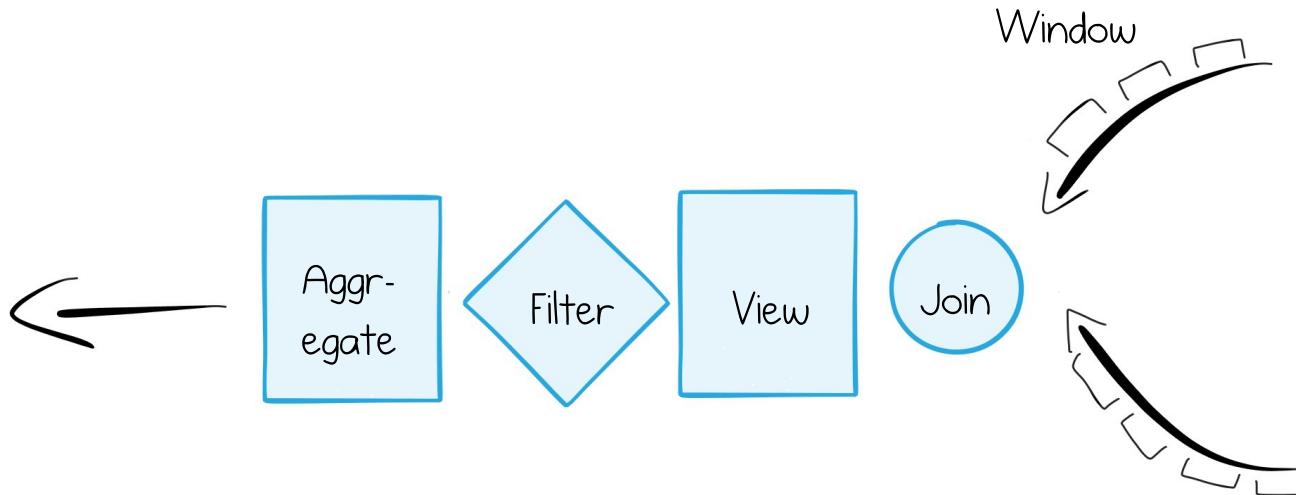
DB Engine designed to process streams

```
Max(price)  
From orders  
where ccy='GBP'  
over 1 day window  
emitting every second
```



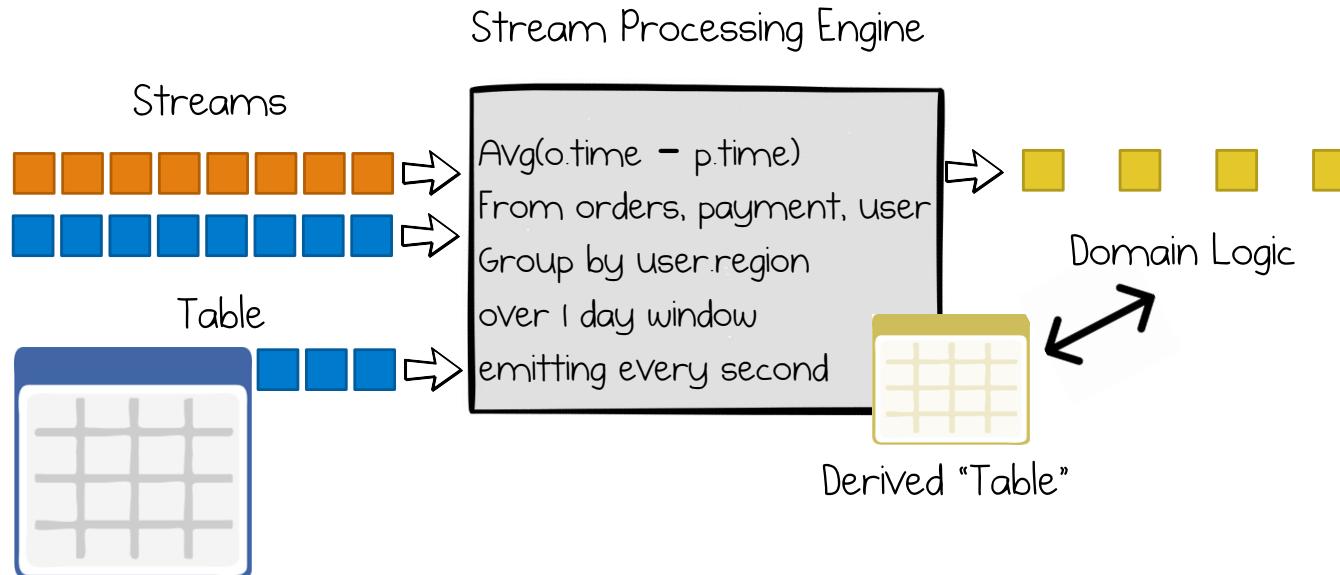


Features: similar to database query engine



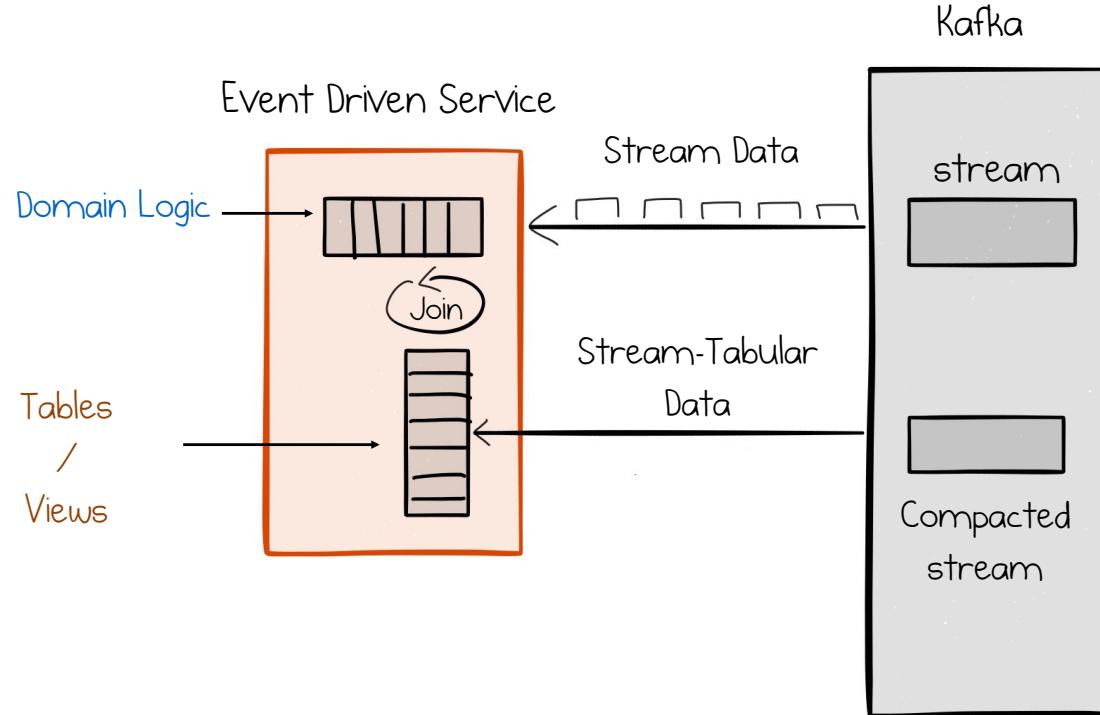


Stateful Stream Processing



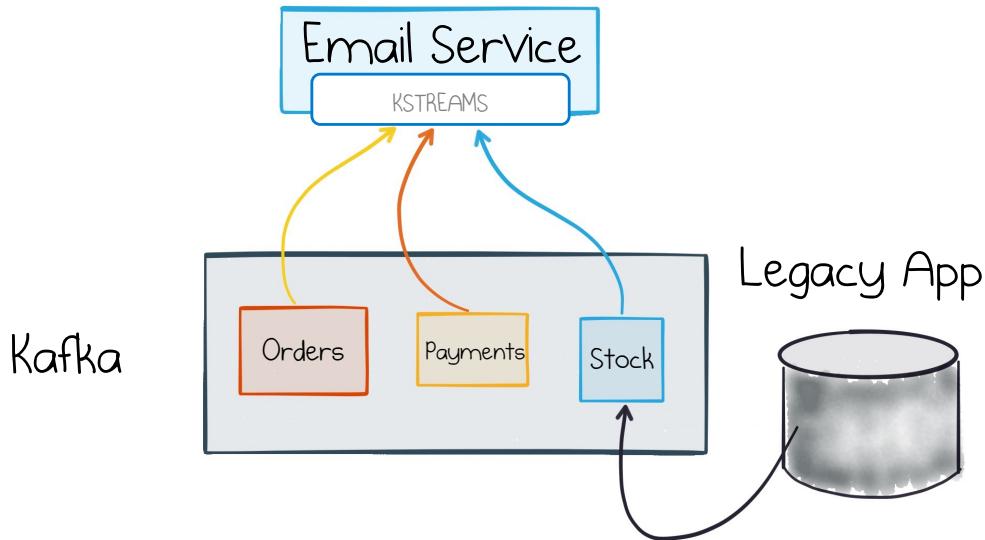


Stateful Stream Processing

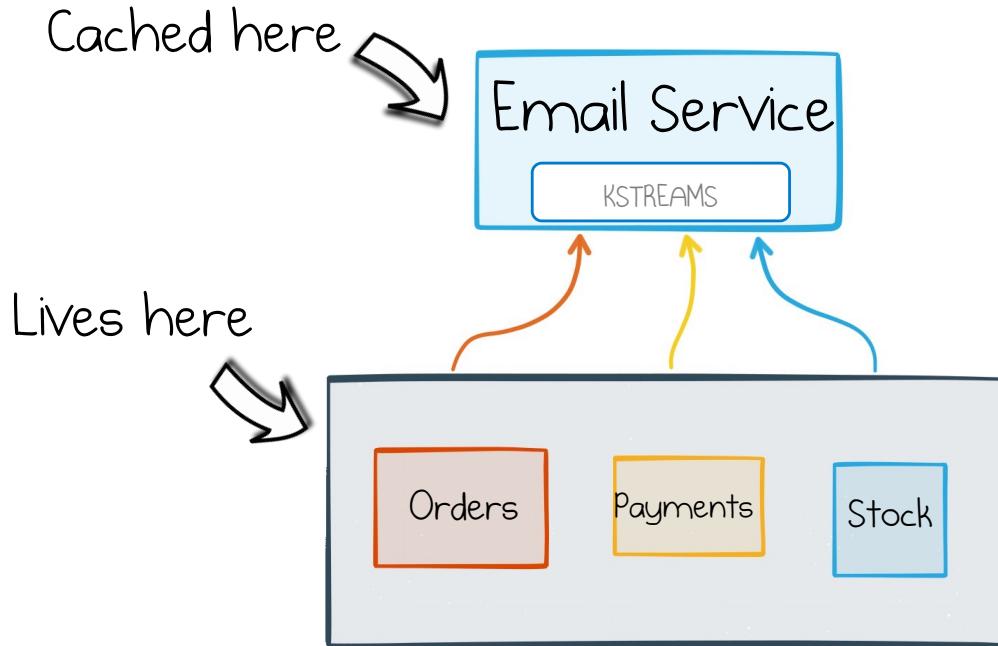




Join shared streams from many other services

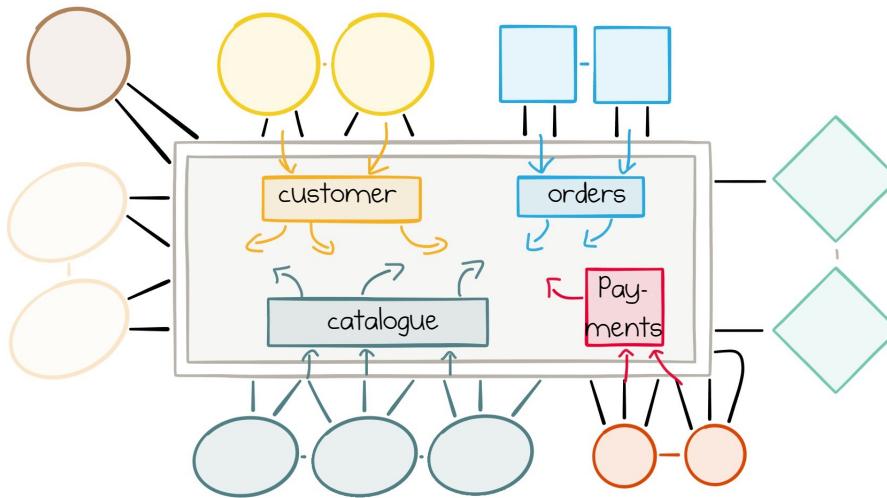


Shared State is only cached in the service,
so there is no way to diverge





Tool for accessing shared, retentive streams



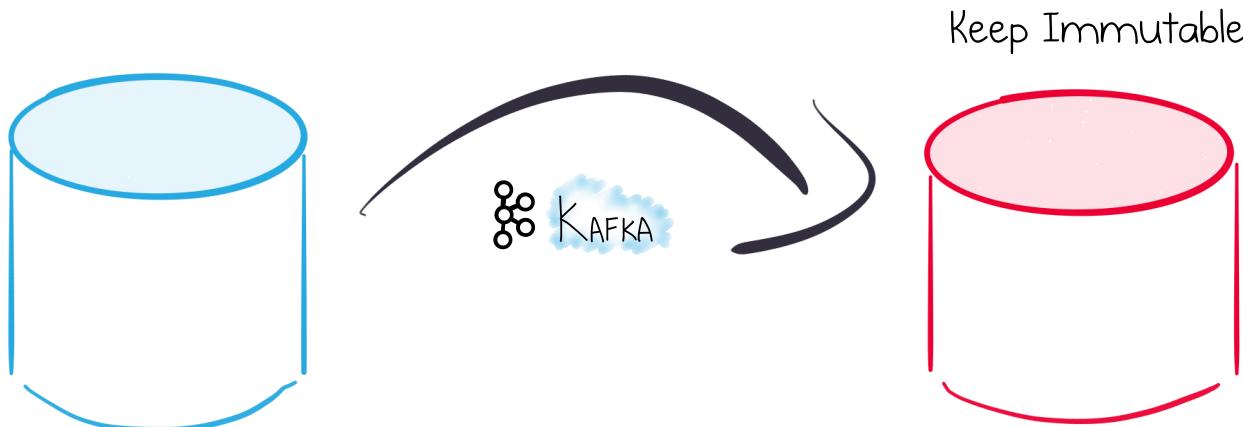


But sometimes you have
to move data



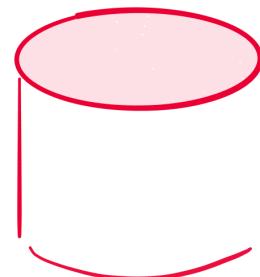


Replicate it, so both copies
are identical



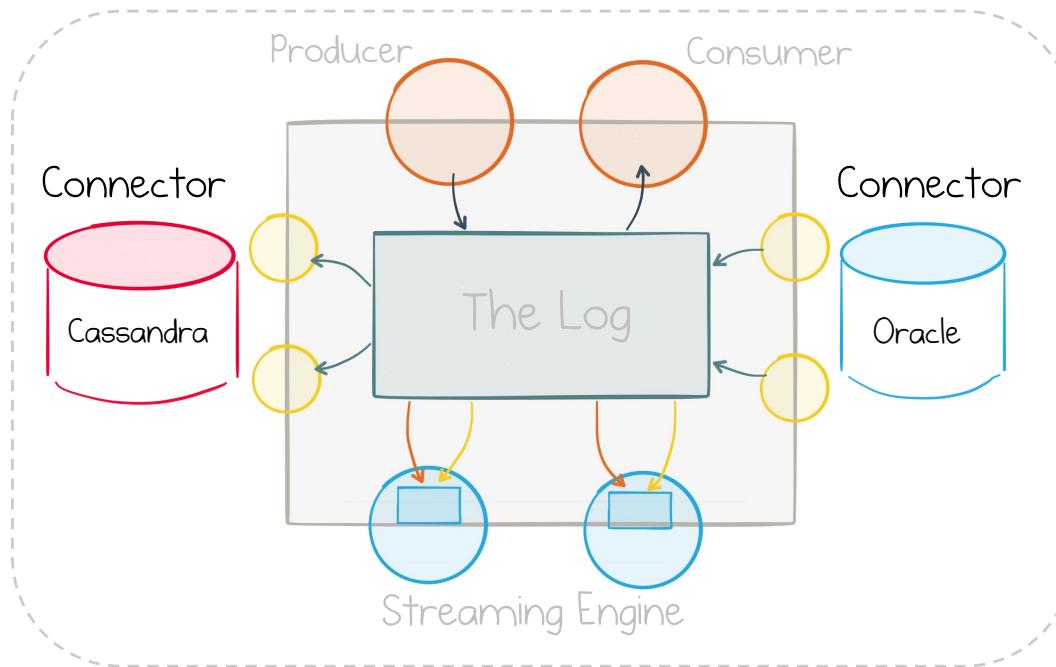


Take only the data you need
today





Confluent's Connectors make this easier





So . . .



When building
services consider
more than just
REST



Remember:

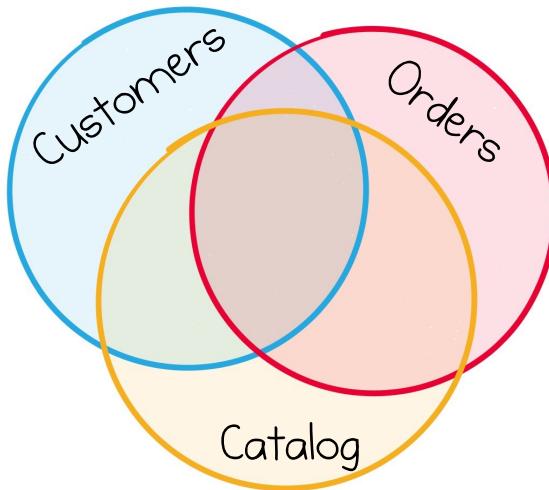
THE DATA DICHOTOMY

Data systems are about exposing data.

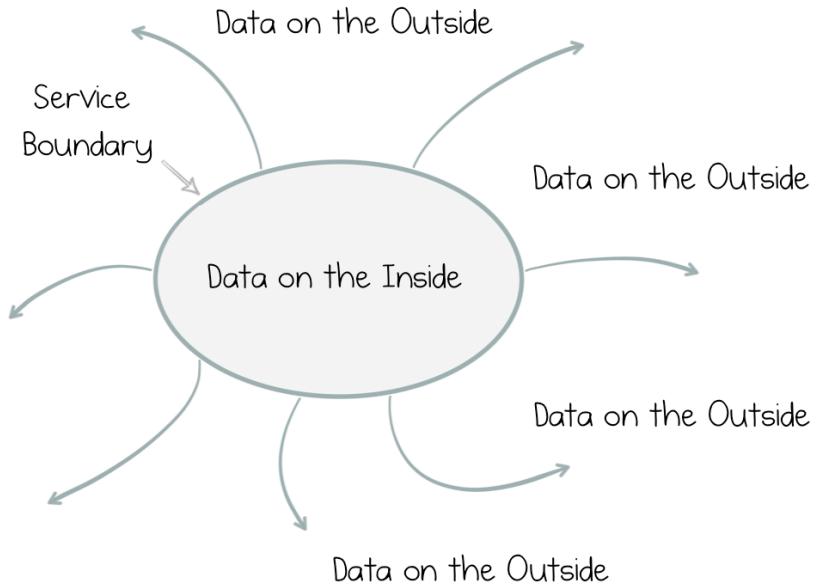
Services are about hiding it.



Shared data is a reality for
most organisations



Embrace the data that both lives and flows between services

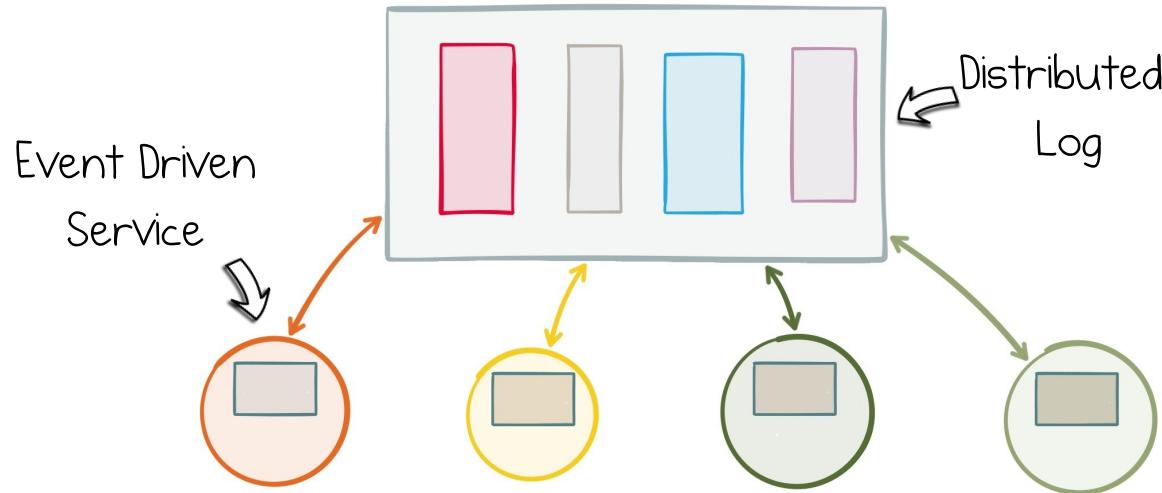




Avoid data services that
have complex / amplifying
interfaces



Share Data via Immutable Streams

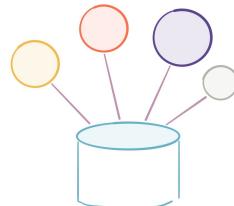


Embed Function into each service

Middle-ground between shared database, messaging and service interfaces



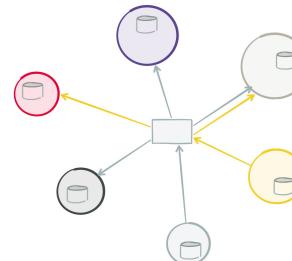
(A) Shared Database



Data & Function
Centralized

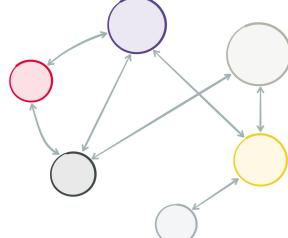
Middle Ground

(B) Messaging



Data & Function
Everywhere

(C) Service Interfaces

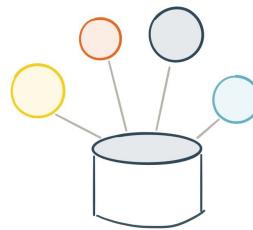


Data & Function in
Owning Services



Balance the Data Dichotomy

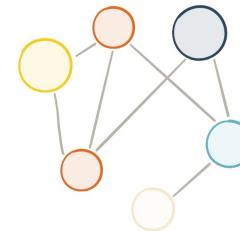
Shared
database



Better
Accessibility



Service
Interfaces



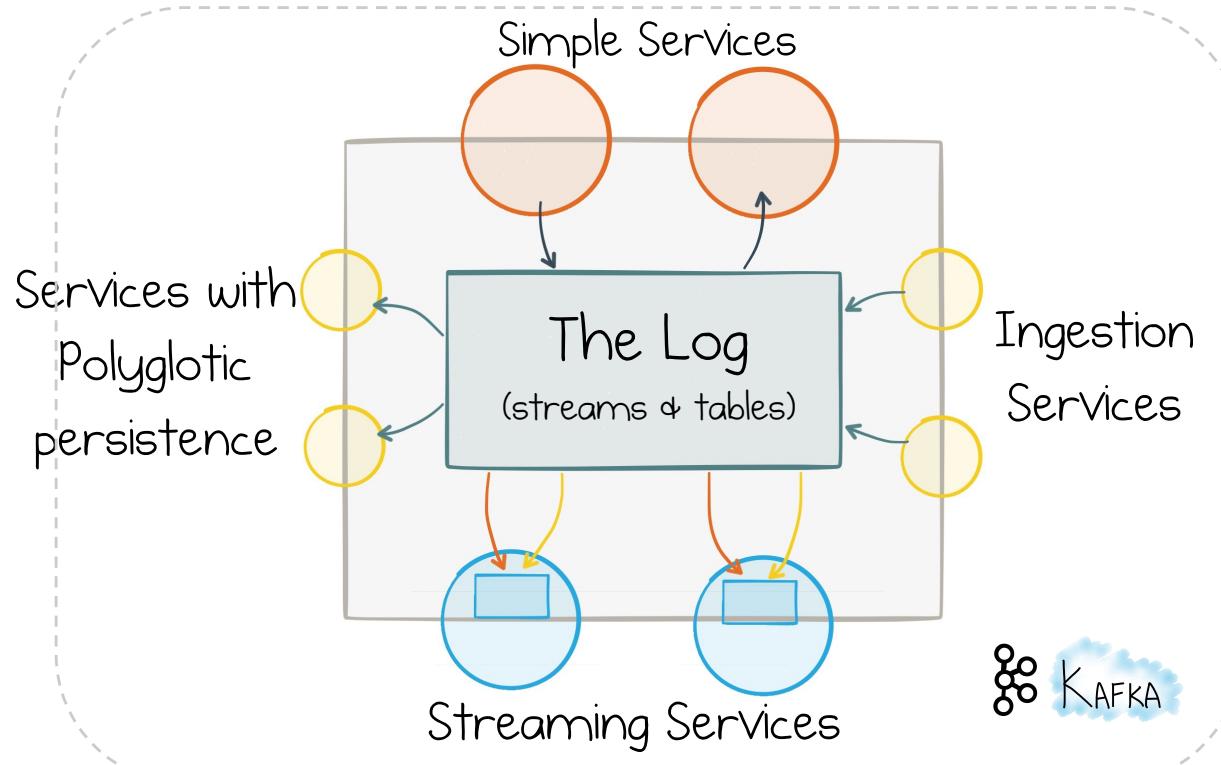
Better
Independence



RIDER Principles

- Reactive - Leverage Asynchronicity. Focus on the now.
- Immutable - Build an immutable, shared narrative.
- Decentralized - No GOD services. Receiver driven.
- Evolutionary - Use only the data you need today.
- Retentive - Rely on the central event log, indefinitely.

Event Driven Service Backbone





Tune in next time

- How do we actually build these things?
- Putting the micro into microservices



Stay in touch!

Online Talks

cnfl.io/online-talks

